

# Bugs as Deviant Behavior: A General Approach to Inferring Errors in Systems Code

Dawson Engle, David Yu Chen, Seth Hallem,  
Andy Chou, Benjamin Chelf

Miroslav Cupak

10/01/2012

- systems obey rules for correctness and performance
- verification, testing, manual and automatic inspection
- problem with static analysis: what rules to check?
- solution: automatically extract from source

- problem: finding what is incorrect without knowing what is correct
- tools: contradictions & common behaviour
- approach: internal consistency & statistical analysis
- implement checkers and apply them to Linux & OpenBSD

- beliefs
  - MUST
  - MAY
- templates
- slots

- checkers defined by:
  - the rule template  $T$
  - the valid slot instances for  $T$
  - the code actions implying beliefs
  - the rules for how beliefs combine & the rules for contradiction
  - the rules for belief propagation
- issues:
  - inferring beliefs (direct observation, pre/post-conditions)
  - relating code (implementation, abstraction)

- internal consistency checkers with modifications:
  - assume all slot-instance combinations are MUST beliefs
  - indicate checks and failures
  - rank the errors and order the results to get most relevant results
- issues:
  - large set of cases
  - pre-processing
  - noise (large samples, ranking, human-level operations)
- extensible, makes use of empty templates
- latent specifications (naming conventions, error codes)

- static null pointer detection
- based on internal consistency
- associates pointers with belief sets and flags contradictions
  - check-then-use
  - use-then-check
  - redundant checks

```
/* 2.4.1:drivers/isdn/avmbi/capidrv.c: */  
1: if (card == NULL) {  
2:     printk(KERN_ERR "capidrv-%d: ... %d!\n",  
3:             card->contrnr, id);  
4: }
```

- detection of accesses to shared variables without their locks
- problem: finding variable-lock bindings
- test MAY beliefs, rank errors
- forward and backward propagation of locks



- determine whether the pointer is a kernel (safe) or a user (tainted) pointer, report intersection
- based on dereference counts
- problems:
  - false positives problem due to kernel backdoors  
checking if they are called from user or kernel code
  - manual inspection

- find routines that are not checked or are incorrectly checked for failures
  - routines returning null pointers are checked before use
  - unnecessary checking of routines that cannot fail

- no A after B (free memory)
- B must follow A (lock-unlock)
- important preprocessing of traces

- finding bugs without special knowledge of the correctness of the program
- many interesting bugs
- new surprising bugs
- real bug reports and patches
- significant portion of false positives (bug/false positives ratio: null pointer 205/40, security 35/19)

- extension of Checking System Rules Using System-Specific, Programmer-Written Compiler Extensions (Using Meta-level Compilation to Check FLASH Protocol Code)
- type systems
- specifications
- dynamic invariant inference

Questions?

What else can we check? Any ideas for other checkers or situations?

## What else can we check?

- check permissions before writing to data structures
- reenable interrupts after disabling them
- size limit on variables
- hold read locks if variable is not modified
- memory allocation before its use
- protocol headers
- double if statements
- waiting on synchronous sends
- certain references not allowed in parts of code



How can we use this tool to prevent deadlocks?

## How can we use this tool to prevent deadlocks?

- lock releasing after acquiring it
- kernel cannot call blocking functions with interrupts disabled
- thread holding spin lock cannot block
- temporal ordering

What do you see as the biggest problem of the proposed tools? How would you tackle it?

# What do you see as the biggest problem of the proposed tools?

- setup cost
- manual inspection
- false positives (ranking, thresholds)
- performance

If you were a developer in a real SW company trying to extend their QE processes, would you consider adoption of this tool a good idea? What could make it more practical?

## Do you consider the tool usable in real SW companies?

- still a lot of manual work, but it can help
- relatively easily extensible
- flexibility
- scalability
- + IDE integration
- + improve performance (storing history, running on parts of code, ignore certain paths)

Are the metrics in use objective?

## Are the metrics in use objective?

- exhaustive all-paths search shifting the rating
- preprocesssing step
- heuristics
- relying on code convention