Network Statistics for OpenFlow

Miroslav Cupák, Masoud Hosseinkhani, Maryam Samizadeh

05/12/2012

Introduction

- OpenFlow provides some very basic statistics per flow
- Network Management requires wider statistics to make wise decisions about networks
- New metrics are needed to be measured

Floodlight Controller

- Our group decided to work on the Floodlight controller
- Floodlight is a Java-based controller
- Widely used and supported by the SDN community

Initial Thoughts!

- Our focus at first was on latency and jitter
- Data-plane metrics!
- OpenFlow is really a control-plane protocol
- How accurate they can be!?
- Let's work on new useful metrics!

Architecture

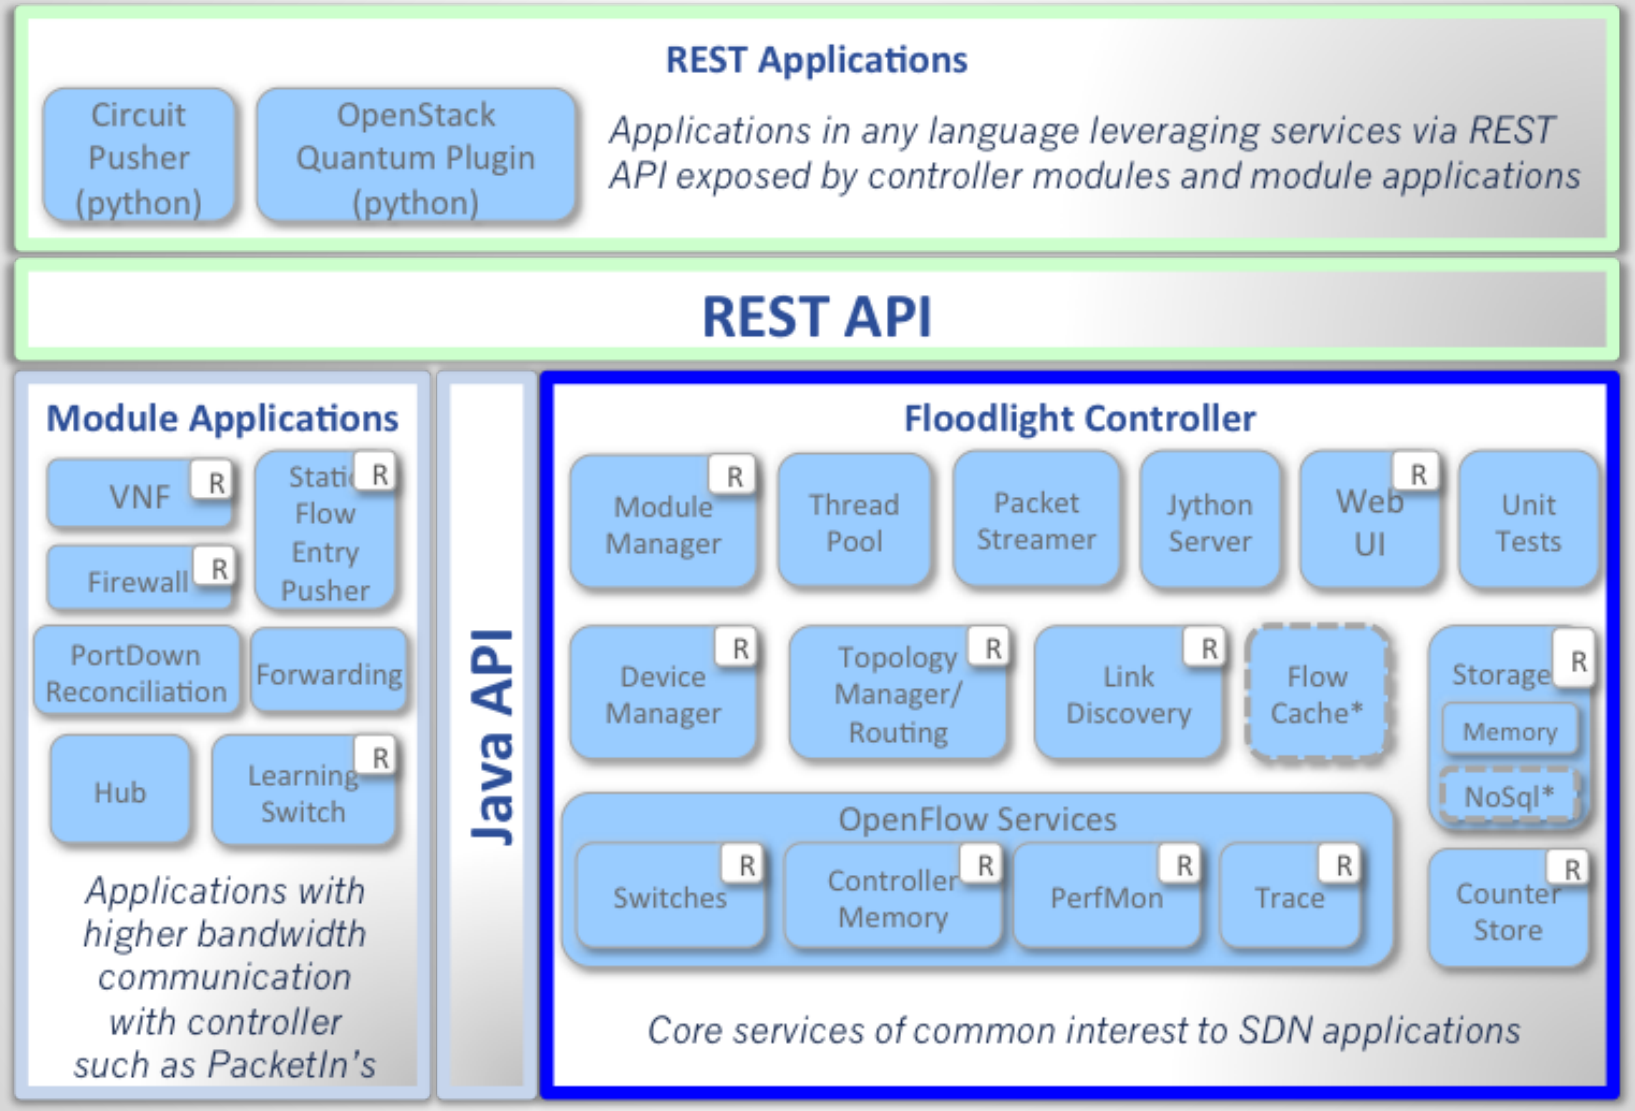

* Interfaces defined only & not implemented: FlowCache, NoSql

Architecture

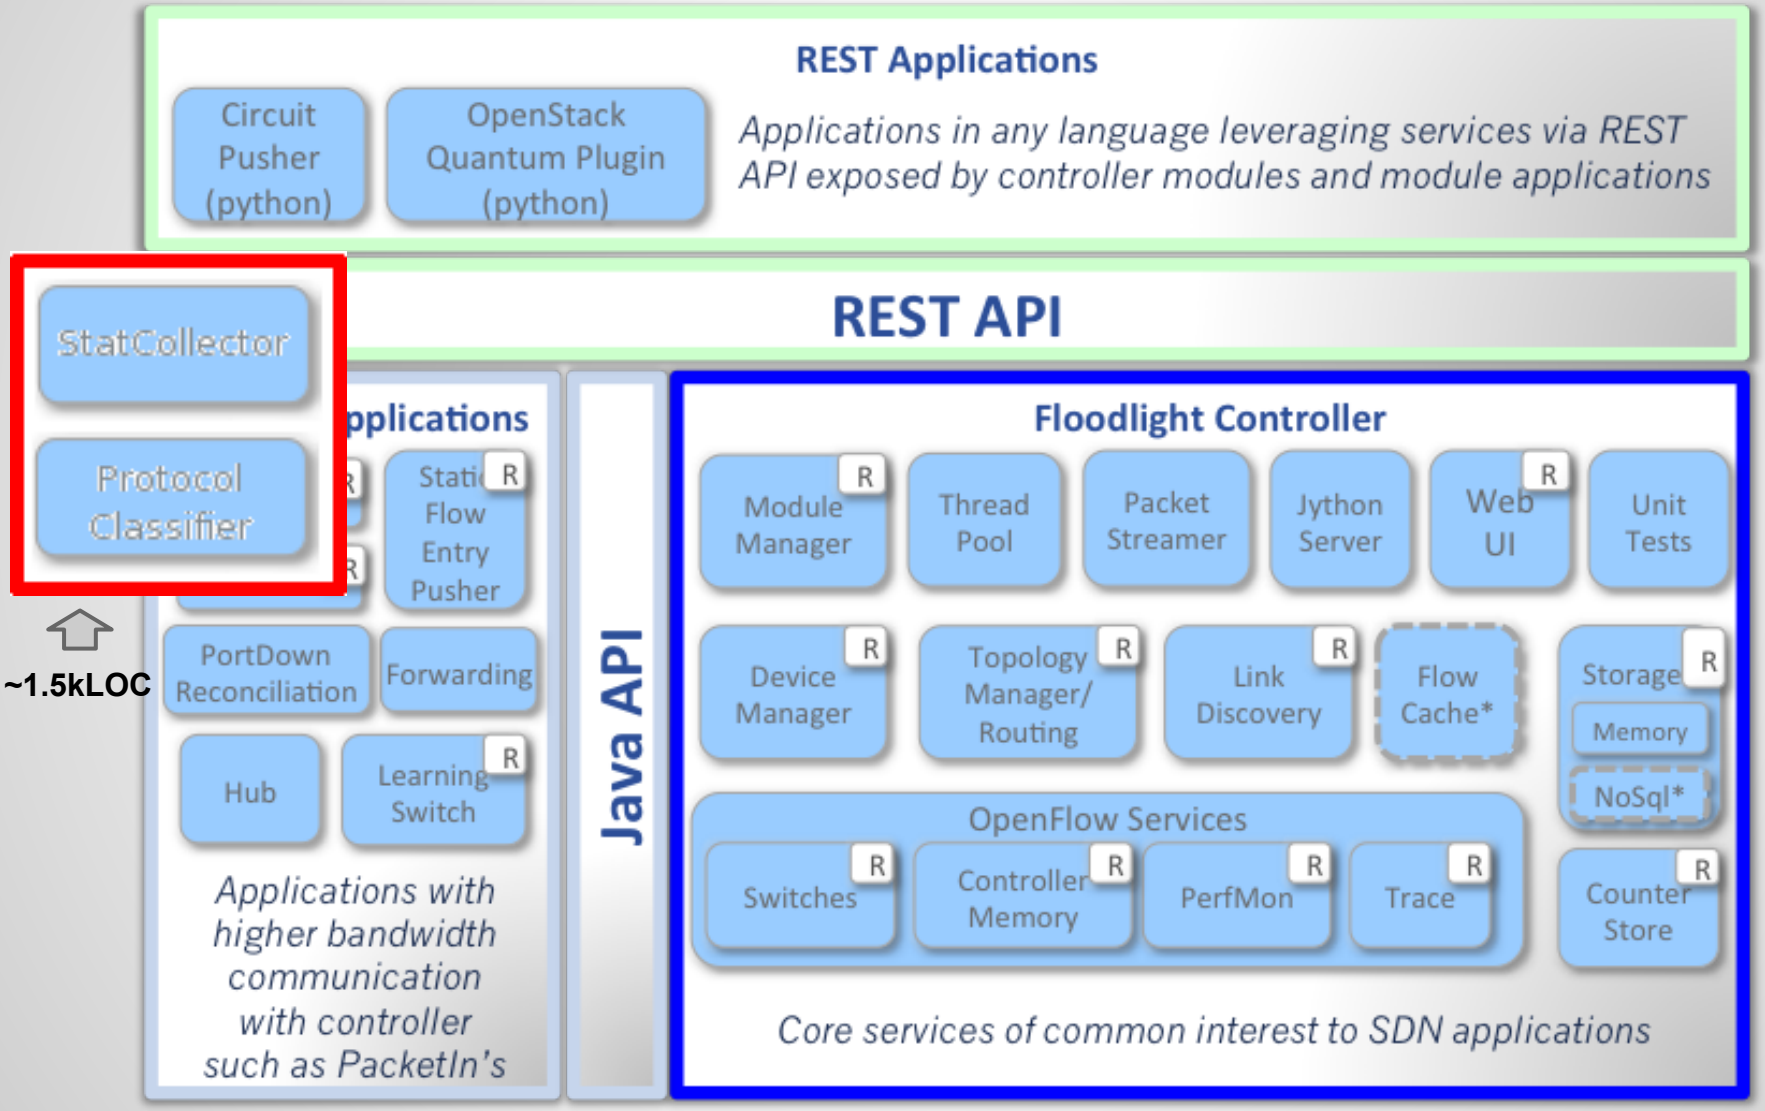

* Interfaces defined only & not implemented: FlowCache, NoSql

Link Bandwidth

- what?
  - consumed data resources, current speed of the links
  - not theoretical maximal throughput, real bandwidth!
- why?
  - planning/scheduling, load balancing, troubleshooting
- how?
  - controller configured to load StatCollector module at startup
  - a periodic task scheduled to run in a separate thread
  - stats for the last interval serialized and available via REST as JSON

Link Bandwidth

- how?
  - finding links - parsing the topology
    - based on LinkDiscoveryManager and TopologyService
    - LLDP packets (a link established if an LLDP is sent out one port and the same LLDP is received on another port)
  - determining real bandwidth
    - based on byte per-port counters
    - periodic queries of switches to get the total value
    - subtract the values measured last time to track values over time
    - average over sent+received data on both ends and divide by time

Link Bandwidth

- how accurate?
  - based on per-port switch counters
  - switches might update the counters only once every couple of seconds
  - limited to switches attached to the controller
  - out of date by the time they're processed by the controller (insignificant)
  - UI displays the last completed interval
  - different delay for switches (reduced by average)

Port Bandwidth

- what?
  - more detailed view on each end of a link
- why?
  - step towards most of the other stats
- how?
  - periodic task started at startup exposing data (REST)
  - no need to search the topology
- how accurate?
  - like link bandwidth except for switch differences

Flow Bandwidth

- what?
  - data over time link usage by flows
- why?
  - proposed metric, interesting insight into network usage
- how?
  - periodic task started at startup exposing data (REST)
  - active polling of switches for flows (duration, B count)
  - average bandwidth over the total duration of the flow
- how accurate?
  - same out-of-date risk as port bandwidth
  - might miss some flows in case of big intervals

Switch Load

- what?
  - number of bytes processed by each switch
- why?
  - network upgrade planning, bottleneck finding
  - computing the list of busiest switches in a network
- how?
  - periodic task started at startup exposing data (REST)
  - aggregating bandwidth data of all the ports of the switch over an interval
- how accurate?
  - same out-of-date risk as port bandwidth

Device Activity

- what?
  - list of the hosts that are most active in the network
- why?
  - useful for network management by providing hints (weak points of the network, hosts with malicious actions)
- how?
  - analyzing incoming packets to find the host creating the flow + its port
  - querying the port for the total number of bytes
  - subtract the values measured last time

Protocol Classification

- what?
  - providing classification of the network traffic based on the higher-level protocols.
- why?
  - knowledge of the type of traffic in the network
  - performance improvements
    - tailoring the structure for more specific traffic.
- how?
  - parsing IP header for the protocol field, mapping the value against the list of L3+ protocols
- how Accurate?
  - very, but only analyzing 1 packet/flow on the controller

Conclusions

- measured metrics that are useful to network management for planning, scheduling, and load balancing
- link/port/flow bandwidth, switch load, device activity, protocol classification
- more focused on the control-plane metrics to be more accurate
- future work: based on the feedback from Floodlight community

Questions?